

```

import pandas as pd
import numpy as np
import random as rnd

import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Activation
from tensorflow.keras.optimizers import Adam

from sklearn.metrics import mean_squared_error, mean_absolute_error, explained_variance_score
from sklearn.metrics import classification_report, confusion_matrix

df=pd.read_csv("/content/kc_house_data.csv")

print(df.columns.values)

['id' 'date' 'price' 'bedrooms' 'bathrooms' 'sqft_living' 'sqft_lot'
 'floors' 'waterfront' 'view' 'condition' 'grade' 'sqft_above'
 'sqft_basement' 'yr_built' 'yr_renovated' 'zipcode' 'lat' 'long'
 'sqft_living15' 'sqft_lot15']

df.head()
df.tail()
df.isnull().sum()
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21613 entries, 0 to 21612
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  -

```

```
0  id          21613 non-null int64
1  date        21613 non-null object
2  price       21613 non-null float64
3  bedrooms   21613 non-null int64
4  bathrooms  21613 non-null float64
5  sqft_living 21613 non-null int64
6  sqft_lot    21613 non-null int64
7  floors      21613 non-null float64
8  waterfront  21613 non-null int64
9  view        21613 non-null int64
10 condition  21613 non-null int64
11 grade       21613 non-null int64
12 sqft_above  21613 non-null int64
13 sqft_basement 21613 non-null int64
14 yr_built    21613 non-null int64
15 yr_renovated 21613 non-null int64
16 zipcode     21613 non-null int64
17 lat         21613 non-null float64
18 long        21613 non-null float64
19 sqft_living15 21613 non-null int64
20 sqft_lot15   21613 non-null int64
dtypes: float64(5), int64(15), object(1)
memory usage: 3.5+ MB
```


```
df.describe().transpose()
```

	count	mean	std	min	25%	50%	75%	max
id	21613.0	4.580302e+09	2.876566e+09	1.000102e+06	2.123049e+09	3.904930e+09	7.308900e+09	9.900000e+09
price	21613.0	5.400881e+05	3.671272e+05	7.500000e+04	3.219500e+05	4.500000e+05	6.450000e+05	7.700000e+06
bedrooms	21613.0	3.370842e+00	9.300618e-01	0.000000e+00	3.000000e+00	3.000000e+00	4.000000e+00	3.300000e+01
bathrooms	21613.0	2.114757e+00	7.701632e-01	0.000000e+00	1.750000e+00	2.250000e+00	2.500000e+00	8.000000e+00
sqft_living	21613.0	2.079900e+03	9.184409e+02	2.900000e+02	1.427000e+03	1.910000e+03	2.550000e+03	1.354000e+04
sqft_lot	21613.0	1.510697e+04	4.142051e+04	5.200000e+02	5.040000e+03	7.618000e+03	1.068800e+04	1.651359e+06
floors	21613.0	1.494309e+00	5.399889e-01	1.000000e+00	1.000000e+00	1.500000e+00	2.000000e+00	3.500000e+00
waterfront	21613.0	7.541757e-03	8.651720e-02	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	1.000000e+00
view	21613.0	2.343034e-01	7.663176e-01	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	4.000000e+00
condition	21613.0	3.409430e+00	6.507430e-01	1.000000e+00	3.000000e+00	3.000000e+00	4.000000e+00	5.000000e+00
grade	21613.0	7.656873e+00	1.175459e+00	1.000000e+00	7.000000e+00	7.000000e+00	8.000000e+00	1.300000e+01
sqft_above	21613.0	1.788391e+03	8.280910e+02	2.900000e+02	1.190000e+03	1.560000e+03	2.210000e+03	9.410000e+03
sqft_basement	21613.0	2.915090e+02	4.425750e+02	0.000000e+00	0.000000e+00	0.000000e+00	5.600000e+02	4.820000e+03
yr_built	21613.0	1.971005e+03	2.937341e+01	1.900000e+03	1.951000e+03	1.975000e+03	1.997000e+03	2.015000e+03
yr_renovated	21613.0	8.440226e+01	4.016792e+02	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	2.015000e+03
zipcode	21613.0	9.807794e+04	5.350503e+01	9.800100e+04	9.803300e+04	9.806500e+04	9.811800e+04	9.819900e+04
lat	21613.0	4.756005e+01	1.385637e-01	4.715590e+01	4.747100e+01	4.757180e+01	4.767800e+01	4.777760e+01
long	21613.0	-1.222139e+02	1.408283e-01	-1.225190e+02	-1.223280e+02	-1.222300e+02	-1.221250e+02	-1.213150e+02
sqft_living15	21613.0	1.986552e+03	6.853913e+02	3.990000e+02	1.490000e+03	1.840000e+03	2.360000e+03	6.210000e+03
sqft_lot15	21613.0	1.276846e+04	2.730418e+04	6.510000e+02	5.100000e+03	7.620000e+03	1.008300e+04	8.712000e+05

```
sns.set(style="whitegrid", font_scale=1)
```

```
plt.figure(figsize=(13,13))
plt.title('Pearson Correlation Matrix',fontsize=25)
sns.heatmap(df.corr(),linewidths=0.25,vmax=0.7,square=True,cmap="GnBu",linecolor='w',
            annot=True, annot_kws={"size":7}, cbar_kws={"shrink": .7})
```

```
<ipython-input-21-1b9cab547edf>:5: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future
sns.heatmap(df.corr(),linewidths=0.25,vmax=0.7,square=True,cmap="GnBu",linecolor='w',
<Axes: title={'center': 'Pearson Correlation Matrix'}>
```



bedr
bathr
sqft_l
sq
1
water
conc
c
sqft_a
sqft_base
yr_
yr_renov
zip
sqft_livi
sqft_

```
price_corr = df.corr()['price'].sort_values(ascending=False)
print(price_corr)
```

```
<ipython-input-22-7cbd0902dff7>:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to
price_corr = df.corr()['price'].sort_values(ascending=False)
```

```
price          1.000000
sqft_living    0.702035
grade          0.667434
sqft_above     0.605567
sqft_living15  0.585379
bathrooms      0.525138
view           0.397293
sqft_basement  0.323816
bedrooms       0.308350
lat            0.307003
waterfront     0.266369
floors         0.256794
yr_renovated   0.126434
sqft_lot       0.089661
sqft_lot15     0.082447
yr_built       0.054012
condition      0.036362
long           0.021626
id            -0.016762
zipcode        -0.053203
Name: price, dtype: float64
```

```
f, axes = plt.subplots(1, 2, figsize=(15,5))
sns.distplot(df['price'], ax=axes[0])
sns.scatterplot(x='price', y='sqft_living', data=df, ax=axes[1])
sns.despine(bottom=True, left=True)
axes[0].set(xlabel='Price in millions [USD]', ylabel='', title='Price Distribution')
axes[1].set(xlabel='Price', ylabel='Sqft Living', title='Price vs Sqft Living')
axes[1].yaxis.set_label_position("right")
axes[1].yaxis.tick_right()
```


<ipython-input-23-8992c7a9a438>:2: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see



<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>



```
sns.set(style="whitegrid", font_scale=1)
```

```
f, axes = plt.subplots(1, 2, figsize=(15,5))
sns.boxplot(x=df['bedrooms'], y=df['price'], ax=axes[0])
sns.boxplot(x=df['floors'], y=df['price'], ax=axes[1])
sns.despine(bottom=True, left=True)
axes[0].set(xlabel='Bedrooms', ylabel='Price', title='Bedrooms vs Price Box Plot')
axes[1].set(xlabel='Floors', ylabel='Price', title='Floors vs Price Box Plot')
```

```
[Text(0.5, 0, 'Floors'),
Text(0, 0.5, 'Price'),
Text(0.5, 1.0, 'Floors vs Price Box Plot')]
```



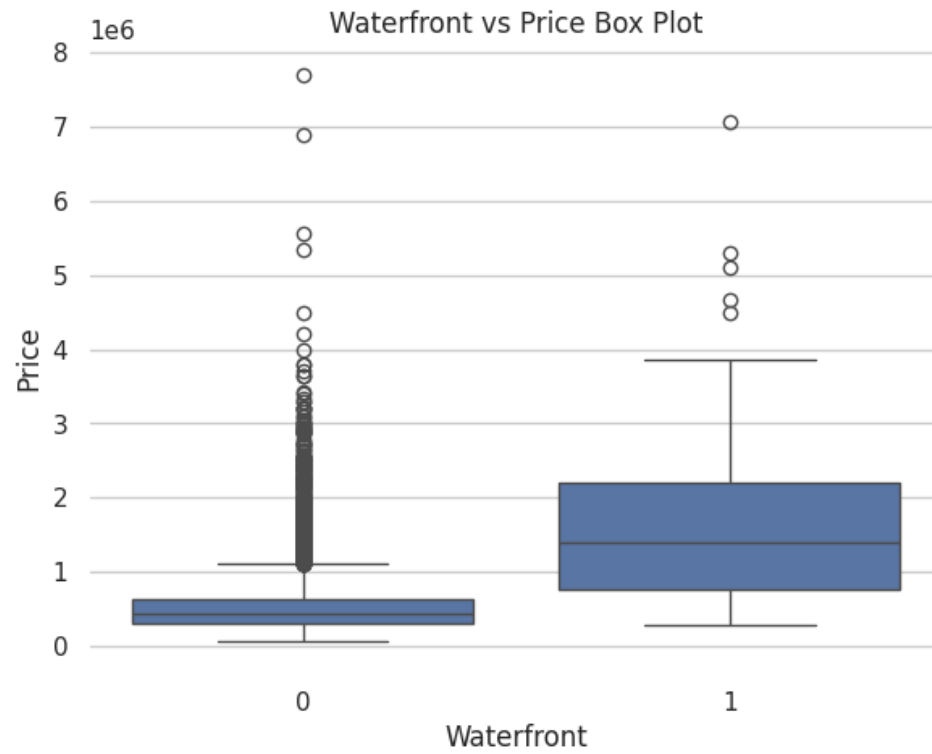



```
f, axes = plt.subplots(1, 2,figsize=(15,5))
sns.boxplot(x=df['waterfront'],y=df['price'], ax=axes[0])
sns.boxplot(x=df['view'],y=df['price'], ax=axes[1])
sns.despine(left=True, bottom=True)
axes[0].set(xlabel='Waterfront', ylabel='Price', title='Waterfront vs Price Box Plot')
axes[1].set(xlabel='View', ylabel='Price', title='View vs Price Box Plot')
```

```
f, axe = plt.subplots(1, 1,figsize=(15,5))
sns.boxplot(x=df['grade'],y=df['price'], ax=axe)
sns.despine(left=True, bottom=True)
axe.set(xlabel='Grade', ylabel='Price', title='Grade vs Price Box Plot')
```



```
[Text(0.5, 0, 'Grade'),
Text(0, 0.5, 'Price'),
Text(0.5, 1.0, 'Grade vs Price Box Plot')]
```





```
df = df.drop('id', axis=1)
df = df.drop('zipcode',axis=1)
```

```
df['date'] = pd.to_datetime(df['date'])
```

```
df['month'] = df['date'].apply(lambda date:date.month)
df['year'] = df['date'].apply(lambda date:date.year)
```

```
df = df.drop('date',axis=1)
```

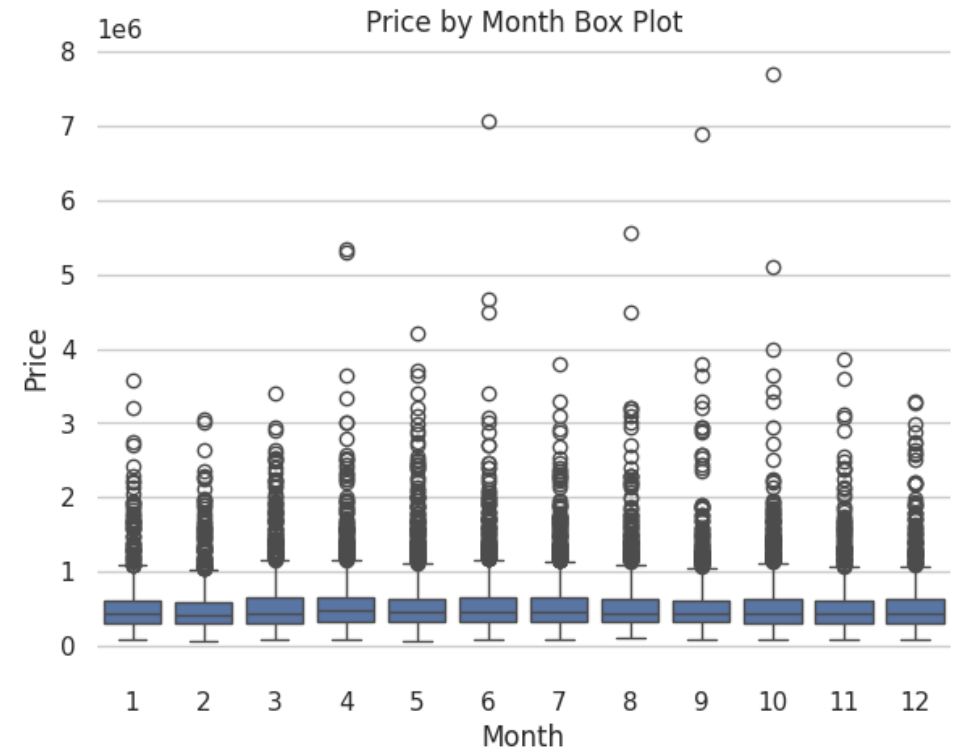

```
# Check the new columns
print(df.columns.values)
```


```
['price' 'bedrooms' 'bathrooms' 'sqft_living' 'sqft_lot' 'floors'
 'waterfront' 'view' 'condition' 'grade' 'sqft_above' 'sqft_basement'
 'yr_built' 'yr_renovated' 'lat' 'long' 'sqft_living15' 'sqft_lot15'
 'month' 'year']
```

```
f, axes = plt.subplots(1, 2,figsize=(15,5))
sns.boxplot(x='year',y='price',data=df, ax=axes[0])
sns.boxplot(x='month',y='price',data=df, ax=axes[1])
sns.despine(left=True, bottom=True)
axes[0].set(xlabel='Year', ylabel='Price', title='Price by Year Box Plot')
axes[1].set(xlabel='Month', ylabel='Price', title='Price by Month Box Plot')
```

```
f, axe = plt.subplots(1, 1,figsize=(15,5))
df.groupby('month').mean()['price'].plot()
sns.despine(left=True, bottom=True)
axe.set(xlabel='Month', ylabel='Price', title='Price Trends')
```

```
[Text(0.5, 0, 'Month'), Text(0, 0.5, 'Price'), Text(0.5, 1.0, 'Price Trends')]
```





```
# Features
X = df.drop('price',axis=1)

# Label
y = df['price']

# Split
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.3,random_state=101)

print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)

(15129, 19)
(6484, 19)
(15129,)
(6484,)

scaler = MinMaxScaler()

# fit and transform
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# everything has been scaled between 1 and 0
print('Max: ',X_train.max())
print('Min: ', X_train.min())

Max:  1.0000000000000002
Min:  0.0
```

```
model = Sequential()

# input layer
model.add(Dense(19,activation='relu'))

# hidden layers
model.add(Dense(19,activation='relu'))
model.add(Dense(19,activation='relu'))
model.add(Dense(19,activation='relu'))

# output layer
model.add(Dense(1))

model.compile(optimizer='adam',loss='mse')

model.fit(x=X_train,y=y_train.values,
          validation_data=(X_test,y_test.values),
          batch_size=128,epochs=400)
```

```

119/119 [=====] - 0s 4ms/step - loss: 28519084032.0000 - val_loss: 28249851904.0000
Epoch 387/400
119/119 [=====] - 0s 3ms/step - loss: 28461717504.0000 - val_loss: 28227166208.0000
Epoch 388/400
119/119 [=====] - 0s 3ms/step - loss: 28447885312.0000 - val_loss: 28236111872.0000
Epoch 389/400
119/119 [=====] - 0s 4ms/step - loss: 28458059776.0000 - val_loss: 28177870848.0000
Epoch 390/400
119/119 [=====] - 0s 4ms/step - loss: 28454924288.0000 - val_loss: 28218097664.0000
Epoch 391/400
119/119 [=====] - 0s 3ms/step - loss: 28471775232.0000 - val_loss: 28149301248.0000
Epoch 392/400
119/119 [=====] - 0s 3ms/step - loss: 28431503360.0000 - val_loss: 28182046720.0000
Epoch 393/400
119/119 [=====] - 0s 4ms/step - loss: 28421937152.0000 - val_loss: 28150104064.0000
Epoch 394/400
119/119 [=====] - 0s 4ms/step - loss: 28381530112.0000 - val_loss: 28120223744.0000
Epoch 395/400
119/119 [=====] - 0s 4ms/step - loss: 28387883008.0000 - val_loss: 28121425920.0000
Epoch 396/400
119/119 [=====] - 0s 4ms/step - loss: 28371150848.0000 - val_loss: 28178825216.0000
Epoch 397/400
119/119 [=====] - 0s 3ms/step - loss: 28336060416.0000 - val_loss: 28081702912.0000
Epoch 398/400
119/119 [=====] - 0s 4ms/step - loss: 28331995136.0000 - val_loss: 28110159872.0000
Epoch 399/400
119/119 [=====] - 1s 5ms/step - loss: 28299653120.0000 - val_loss: 28048721920.0000
Epoch 400/400
119/119 [=====] - 1s 5ms/step - loss: 28302182400.0000 - val_loss: 28045254656.0000
<keras.src.callbacks.Historv at 0x7dd02f5f89d0>


```

```
losses = pd.DataFrame(model.history.history)
```

```

plt.figure(figsize=(15,5))
sns.lineplot(data=losses,lw=3)
plt.xlabel('Epochs')
plt.ylabel('')
plt.title('Training Loss per Epoch')
sns.despine()

```



```
# predictions on the test set
predictions = model.predict(X_test)

print('MAE: ',mean_absolute_error(y_test,predictions))
print('MSE: ',mean_squared_error(y_test,predictions))
print('RMSE: ',np.sqrt(mean_squared_error(y_test,predictions)))
print('Variance Regression Score: ',explained_variance_score(y_test,predictions))

print('\n\nDescriptive Statistics:\n',df['price'].describe())
```

```
203/203 [=====] - 2s 7ms/step
MAE: 104748.00371045554
MSE: 28045253475.152184
RMSE: 167467.17133561487
Variance Regression Score: 0.8000695603775322
```

Descriptive Statistics:

```

count    2.161300e+04
mean     5.400881e+05
std      3.671272e+05
min      7.500000e+04
25%      3.219500e+05
50%      4.500000e+05
75%      6.450000e+05
max      7.700000e+06
Name: price, dtype: float64

```

```
f, axes = plt.subplots(1, 2, figsize=(15,5))
```

```

# Our model predictions
plt.scatter(y_test, predictions)

```

```

# Perfect predictions
plt.plot(y_test, y_test, 'r')

```

```

errors = y_test.values.reshape(6484, 1) - predictions
sns.distplot(errors, ax=axes[0])

```

```

sns.despine(left=True, bottom=True)
axes[0].set(xlabel='Error', ylabel='', title='Error Histogram')
axes[1].set(xlabel='Test True Y', ylabel='Model Predictions', title='Model Predictions vs Perfect Fit')

```

```
<ipython-input-36-83b90cb0bedd>:10: UserWarning:
```

```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

```
~\AppData\Local\Programs\Python\Python310\python.exe
```